

Dynamic Demand Prediction for Sustainable Bike Sharing Systems

By Alexander Busch and Kaifeng Lu as part of the course project of ECE1724H: Bio-inspired Algorithms for Smart Mobility. Dr. Alaa Khamis, University of Toronto, 2024.

Literature Review

Scope and Overview

Due to the vastly growing number of bike sharing systems, bike sharing demand prediction has been investigated by a number of authors in the last 15 years. While classical machine learning approaches such as regression and boosting have been used in the early developments, in recent years, deep learning approaches have been found to give significant advantages in the forecasting domain [1].

Based on the literature, we compare the following machine learning approaches to bike sharing demand prediction and identify open problems.

Paper Summaries

Bike-Sharing Demand Prediction at Community Level under COVID-19 Using Deep Learning: [2]

This article investigates short-term bike-sharing demand forecasting in Montreal. It compares deep learning techniques for predicting bike pickups 15 minutes ahead in six communities identified within the city's bike-sharing network. The study compares the performance of LSTMs (Long Short Term Memory, a recurrent network type), CNN (Convolutional Neural Network)-LSTM hybrids, and ARIMA (Auto-Regressive Integrated Moving Average). The authors trained the models based on two main feature attributes: historical demand data in 15-minute intervals and weather conditions.

They found that the CNN-LSTM hybrid model outperforms the other models (MAE of 3.00 and RMSE of 4.77), compared to the LSTM model (MAE of 4.48 and RMSE of 6.86), and the ARIMA (MAE of 50.95 and RMSE of 61.17).

Pros:

- LSTM Models: These were found to capture temporal dependencies effectively with moderate complexity.
- CNN-LSTM Hybrid Models: These were found to combine spatial and temporal learning for improved accuracy in complex data.
- The ARIMA Model is fast to train and easy to interpret.

Cons:

- LSTM Models: Lacks spatial awareness, limiting prediction accuracy compared to hybrid models.
- CNN-LSTM Hybrid Models: These are computationally expensive and harder to interpret.
- The ARIMA Model was found to perform poorly for complex, nonlinear and disrupted data.

2. Using graph structural information about flows to enhance short-term demand prediction in bike-sharing systems [1]:

This paper investigates how graph theory can enhance short-term bike-sharing demand forecasting. It focuses on incorporating graph-based features derived from flow interactions such as Out-strength, In-strength, Out-degree, In-degree, and PageRank to improve prediction accuracy. It

compares three machine learning models: XGBoost, Multi-Layer Perceptron (MLP), and LSTM. This study found that including graph-based features significantly improves model performance compared to using traditional features like meteorological data alone.

The authors here find that the XGBoost, MLP and LSTM model achieve MAPE values of 27.2%, 27.7%, and 27.0% as well as RMSE values of 6.78, 6.85, and 6.69 correspondingly. The LSTM model was here found to be the most effective model for incorporating complex graph-based features, while the overall performance of the XGBoost is nearly comparable.

Pros:

- XGBoost excels at handling structured data and provides good performance with less tuning.
- The MLP architecture is simple and efficient with moderate performance.
- The LSTM is best at capturing temporal dependencies and sequential patterns, especially with time-lagged data.

Cons:

- The XGBoost is less effective at handling complex time dependencies compared to the LSTM.
- The MLP struggles with sequential data and lacks the advanced handling of time dependencies.
- The LSTM is complex to train and requires significant computational resources.

3. Modeling bike-sharing demand using a regression model with spatially varying coefficients [3]:

This article focuses on investigating how various factors such as land use, socio-demographic attributes, and transportation infrastructure influence bike-sharing demand at different stations. Notably, they define a graph model, where the data for each bike-sharing cluster is accumulated according to its catchment region. The catchment region is calculated using Thiessen polygons, and ensured to be non-overlapping. The authors propose a spatially varying coefficients (SVC) regression model that accounts for local spatial effects, unlike previous regression models that assume the factors are spatially homogeneous.

Pros:

- The graph model includes connectivity between station clusters to be used.
- The more sophisticated modeling of station clusters defined by catchment areas instead of single stations notably allows to predict demand in regions where no stations are present and thus allows planning.

Cons:

- As the model does not directly regress on immediate historical data as e.g. the LSTM models, it fundamentally does not allow accurate future predictions, but instead better fits general analysis purposes.

Modeling Bike Availability in a Bike-Sharing System Using Machine Learning [4]:

This paper explores predicting bike availability at San Francisco Bay Area Bike Share stations using machine learning algorithms. The authors apply three methods: Random Forest (RF) and Least-Squares Boosting (LSBoost) for univariate regression and Partial Least-Squares Regression (PLSR) for multivariate regression. They found that factors like station neighbors, time of prediction, and weather conditions are significant in predicting bike availability.

The RF, LSBoost and PLSR models achieve MAE of 0.37, 0.58 and 0.6 bikes per station correspondingly, i.e. the Random Forest (RF) model here is found to give the best prediction accuracy.

Pros:

- The RF model is highly accurate with low prediction error and robust to overfitting.
- The LSBoost model is effective for regression tasks with manageable computational complexity.
- The PLSR captures spatial correlations, useful for large networks with interdependent stations.

Cons:

- RF: Requires independent observations. Performance decreases as prediction horizon increases.
- LSBoost: Higher prediction error. Requires proper regularization against overfitting.
- PLSR: Yields the highest prediction error. Less accurate for smaller networks.

Cross-Mode Knowledge Adaptation for Bike Sharing Demand Prediction Using Domain-Adversarial Graph Neural Networks [5]

This recent paper focuses on improving the state of the art in bike sharing demand prediction by integrating additional features from other transport modes, here subway and ride-hailing traffic data, all from New York City. In order to integrate this heterogeneous graph structured data, recurrent CNNs are combined with graph convolutions to yield transport embeddings. The paper also introduces an adversarial training principle to learn these embeddings, such that the embeddings learned are optimized to be indistinguishable. These embeddings are then fed into a multiple GNNs, combined into a single representation, and passed to a final prediction layer.

Pros:

- The architecture developed seems to give a significant improvement to the compared approaches. However, the dataset itself is one of the most often used in bike-share prediction [1], allowing some degree of comparison of results.
- The architecture allows extension to integrate very heterogeneous other modes of transportation significant to bike-sharing that would otherwise be very difficult to integrate.

Cons:

- The architecture and total approach is complicated, requiring extensive effort to adapt to different formats of data from other systems.
- The data, especially from ride-hailing, is likely more noisy than the bike-sharing trip data (due to e.g. only several providers contributing to the ride-hailing data, and the different companies providing data using their own standards)
- It is unclear how much improvement the architecture would give without the additional features.
- Due to the station structure used for the graph, the model cannot predict demand for only places themselves.

Taxonomic Table and Tabular Comparison

We organize the used models into a tree based on the dominant design features for a brief overview of the common investigated models for bikesharing demand prediction. The fundamental division is between two lines of work between “classical” approaches and deep approaches, which split into different lines of works. Further, because the more advanced deep models build on LSTM-approaches to model immediate history, this marks another division in the tree. In order to have concise names for the approaches, we denote every approach by the fundamental architectural characteristic and its paper (e.g. LSTM [2] denotes the LSTM approach in [2]), although this does neglect other important differences in the approaches.

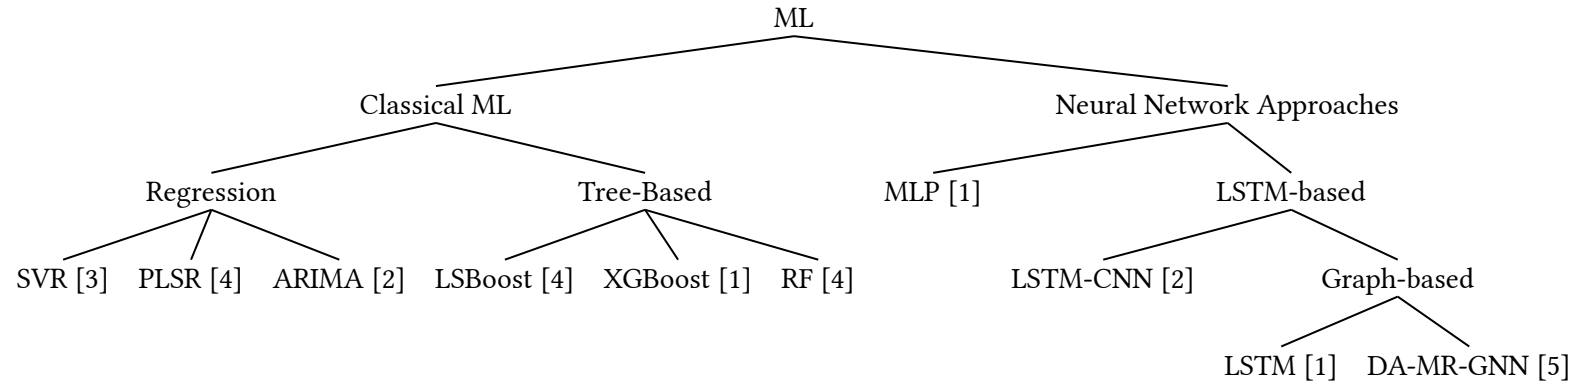


Figure 1: A classification tree of the employed ML approaches

We summarize the main advantages and disadvantages in the following table. We classify how well each approach integrates spatial, recent historical and connectivity information. For example, the pure LSTM approach handles and generalizes recent historical information well, while the regression does not have explicit integration of these features, it only incorporates historical information via engineered features.

Approach	Spatial Handling	Temporal Handling	Connectivity Modeling	Model Complexity	Model Performance
PLSR [4]	Decent	Rudimentary	Rudimentary	Rudimentary	Low
LSBoost in [4]	Rudimentary	Rudimentary	Rudimentary	Low	Decent - Good
RF in [4]	Rudimentary	Rudimentary	Rudimentary	Low	Good
XGBoost in [2]	Rudimentary	Rudimentary	Rudimentary	Low	Good
ARIMA in [2]	Rudimentary	Short Term	Rudimentary	Low	Decent
MLP in [2]	Rudimentary	Short Term	Rudimentary	Rudimentary	Decent
SVR [3]	Good	Rudimentary	Decent - Good	Medium (Complex Feature Engineering)	Good for spatial modeling
DA-MR-GNN [5]	Decent	Very good	Very Good	Very high	Very good

Directions for Future Research

As noted in [1], there is no single standard benchmarking dataset for the specific objective of bike sharing demand prediction, making it difficult to obtain a precise ranking of the models. Notably, the New York Citibike data has been used in multiple approaches [1], [5], but the precise data used still changes. As given in [1], several studies have suggested that XGBoost performing as well most of the state of the art approaches, having also won the 2014 Kaggle competition [6], however the best approach was found to depend strongly on the dataset and modeling.

Notably [1] suggests, that well-performing deep model architectures in traffic prediction are likely to show good performance in bike-sharing demand prediction. In this line of work, [5] combines well performing spatio-temporal GNN approach, known from traffic prediction [7] and [8], with a domain-adversarial network. However, little work has been done on the extent that performance of traffic prediction carries over to bike-sharing performance and adaptation of successful models from other adjacent fields, (e.g. models in [7], [8]), appears to be a promising research direction.

Bibliography

- [1] Y. Yang, A. Heppenstall, A. Turner, and A. Comber, "Using graph structural information about flows to enhance short-term demand prediction in bike-sharing systems," *Computers, Environment and Urban Systems*, vol. 83, p. 101521–101522, 2020, doi: <https://doi.org/10.1016/j.compenvurbsys.2020.101521>.
- [2] A. Mehdizadeh Dastjerdi and C. Morency, "Bike-Sharing Demand Prediction at Community Level under COVID-19 Using Deep Learning," *Sensors*, vol. 22, no. 3, 2022, doi: 10.3390/s22031060.
- [3] X. Wang, Z. Cheng, M. Trépanier, and L. Sun, "Modeling bike-sharing demand using a regression model with spatially varying coefficients," *Journal of Transport Geography*, vol. 93, p. 103059–103060, 2021, doi: <https://doi.org/10.1016/j.jtrangeo.2021.103059>.
- [4] H. I. Ashqar, M. Elhenawy, M. H. Almannaa, A. Ghanem, H. A. Rakha, and L. House, "Modeling bike availability in a bike-sharing system using machine learning," in *2017 5th IEEE International Conference on Models and Technologies for Intelligent Transportation Systems (MT-ITS)*, 2017, pp. 374–378. doi: 10.1109/MTITS.2017.8005700.
- [5] Y. Liang, G. Huang, and Z. Zhao, "Cross-Mode Knowledge Adaptation for Bike Sharing Demand Prediction Using Domain-Adversarial Graph Neural Networks," *IEEE Transactions on Intelligent Transportation Systems*, vol. 25, no. 5, pp. 3642–3653, 2024, doi: 10.1109/TITS.2023.3322717.
- [6] W. Cukierski, "Bike Sharing Demand." Kaggle, 2014.
- [7] T. Yu, H. Yin, and Z. Zhu, "Spatio-temporal Graph Convolutional Neural Network: A Deep Learning Framework for Traffic Forecasting," *ArXiv*, 2017, [Online]. Available: <https://api.semanticscholar.org/CorpusID:4972291>
- [8] S. Zhang, Y. Guo, P. Zhao, C. Zheng, and X. Chen, "A Graph-Based Temporal Attention Framework for Multi-Sensor Traffic Flow Forecasting," *IEEE Transactions on Intelligent Transportation Systems*, vol. 23, pp. 7743–7758, 2022, [Online]. Available: <https://api.semanticscholar.org/CorpusID:234906657>